

Assignment 3: Simplified Agentic RAG System with Gemini & LangGraph

1. Background

Retrieval-Augmented Generation (RAG) combines a knowledge base (KB) with a Large Language Model (LLM) so answers are grounded in factual content. An **agentic RAG** adds a self-review loop: the model critiques its own answer, checks against the KB, and refines if needed.

In this assignment, you will build a **lightweight agentic RAG pipeline** using:

- **Gemini on Google Cloud (Vertex AI)** as the LLM
 - **LangGraph** to wire the workflow
-

2. Problem Statement

“Build a simplified Agentic RAG system with Gemini + LangGraph that retrieves up to 5 KB snippets, generates an initial answer, critiques it, and when necessary refines the answer by pulling one more snippet, returning a citation-backed response.”

Dataset: `self_critique_loop_dataset.json`

Note that the dataset and a pinecone starter file is available over LMS.

3. Detailed Tasks

3.1 Preprocessing & Indexing

1. Load the KB JSON (~30 entries).
 2. Generate embeddings using **Vertex AI Embeddings (models/gemini-embedding-001)**.
 3. Store vectors in a database (Anyone one of: **Pinecone**, **Weaviate**, or **Qdrant**). (recommended is Pinecone)
-

3.2 Define LangGraph Workflow

Your graph should have 4 main nodes:

1. **Retriever Node (retrieve_kb)**

- Input: user question
 - Output: top 5 snippets from vector DB
2. **LLM Answer Node (`generate_answer`)**
 - Model: **Gemini 1.5 (Vertex AI)**
 - Input: question + snippets
 - Output: initial answer with citations [KBxxx]
 3. **Self-Critique Node (`critique_answer`)**
 - Model: Gemini
 - Checks completeness of initial answer vs. snippets
 - Output: "COMPLETE" or "REFINE: <missing keywords>"
 4. **Refinement Node (`refine_answer`)**
 - If refinement needed:
 - Retrieve **1 more snippet** with missing keywords
 - Regenerate answer including this snippet
 - Output: final refined answer

Decision logic:

- If critique = COMPLETE → return initial answer
 - If critique = REFINE → return refined answer
-

4. Tools & Tech

- **LangGraph** (Graph API for agentic flows)
 - **Gemini on GCP (Vertex AI)** for LLM & embeddings
 - **Vector DB:** Pinecone / Weaviate / Qdrant
 - **Python 3.10**
 - Suggested packages:

```
langgraph
google-cloud-aiplatform
pinecone-client    # or weaviate-client / qdrant-client
pydantic
```
-

5. Testing Queries

Try the pipeline with:

1. “*What are best practices for caching?*”

2. *“How should I set up CI/CD pipelines?”*
 3. *“What are performance tuning tips?”*
 4. *“How do I version my APIs?”*
 5. *“What should I consider for error handling?”*
-

6. Deliverables

Submit either:

- **Jupyter Notebook** showing all steps OR
 - **ZIP folder** with:
 - `index_kb.py` (embeddings + vector DB)
 - `agentic_rag_simplified.py` (LangGraph workflow)
 - `requirements.txt`
-

7. Notes

- Keep flow simple (1 critique, 1 refinement at most)
 - Always cite snippets [KBxxx]
 - Use `temperature=0` for consistency
 - Log each step: retrieved snippets, initial answer, critique, final answer using `logger/mlflow`
-